

TIL (Things I Learned) Performance Analysis of CoreCLR Interpreter

Yuchong Pan

Microsoft

August 6, 2020

TIL: What's TIL?

TIL: What's TIL?

- ▶ TIL = Today I Learned

TIL: What's TIL?

- TIL = Today I Learned

[TOP DEFINITION](#)



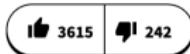
TIL

T.I.L. (Today I Learned)

Often used as a [space saver](#) on websites such as Reddit or [DIGG](#) when writing the titles of links to [interesting things](#) the submitter has learned that day.

TIL that a [meter](#) is [the distance](#) light travels in [299,792,458ths](#) of a second.

by [zBriGuy](#) August 20, 2009



TIL: What's TIL?

- ▶ TIL = Today I Learned

[TOP DEFINITION](#)



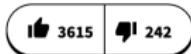
TIL

T.I.L. (Today I Learned)

Often used as a [space saver](#) on websites such as Reddit or [DIGG](#) when writing the titles of links to [interesting things](#) the submitter has learned that day.

TIL that a [meter](#) is [the distance](#) light travels in [299,792,458ths](#) of a second.

by [zBriGuy](#) August 20, 2009



- ▶ TIL = Things I Learned

TIL: What's TIL?

- ▶ TIL = Today I Learned

TOP DEFINITION



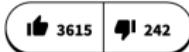
TIL

T.I.L. (Today I Learned)

Often used as a [space saver](#) on websites such as Reddit or [DIGG](#) when writing the titles of links to [interesting things](#) the submitter has learned that day.

TIL that a [meter](#) is [the distance](#) light travels in [299,792,458ths](#) of a second.

by [zBriGuy](#) August 20, 2009



- ▶ TIL = Things I Learned
- ▶ TIL = Things I Learned about the CoreCLR interpreter and its performance during my internship in the summer of 2020

TIL: Who Am I?

TIL: Who Am I?

► Yuchong Pan 潘宇冲

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>
- ▶ UBC Combined Honours Computer Science and Mathematics

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>
- ▶ UBC Combined Honours Computer Science and Mathematics
- ▶ It rains a lot in Vancouver, too.

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>
- ▶ UBC Combined Honours Computer Science and Mathematics
- ▶ It rains a lot in Vancouver, too.
- ▶ Compilers, programming languages, and algorithms!

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>
- ▶ UBC Combined Honours Computer Science and Mathematics
- ▶ It rains a lot in Vancouver, too.
- ▶ Compilers, programming languages, and algorithms!
- ▶ Third internship at Microsoft

TIL: Who Am I?

- ▶ Yuchong Pan 潘宇冲
- ▶ <http://ypan.me>
- ▶ UBC Combined Honours Computer Science and Mathematics
- ▶ It rains a lot in Vancouver, too.
- ▶ Compilers, programming languages, and algorithms!
- ▶ Third internship at Microsoft
- ▶ Last year: .NET Core Uninstall Tool

.NET Core Uninstall Tool

05/27/2020 • 11 minutes to read • 

The [.NET Core Uninstall Tool](#) (`dotnet-core-uninstall`) lets you remove .NET Core SDKs and Runtimes from a system. A collection of options is available to specify which versions you want to uninstall.

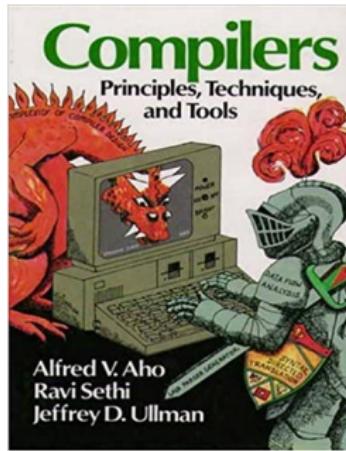
TIL: Who Am I?



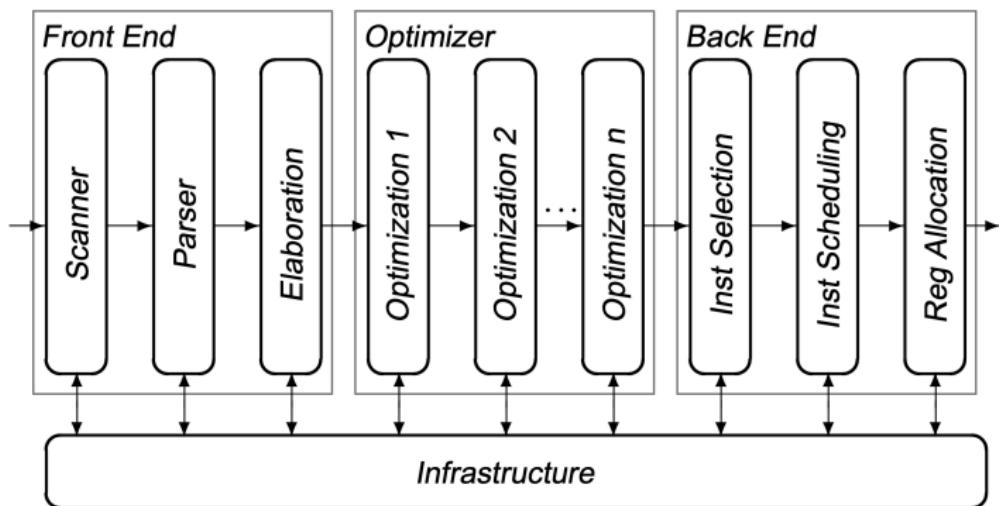
TIL: Why Is the Interpreter So Interesting?

TIL: Why Is the Interpreter So Interesting?

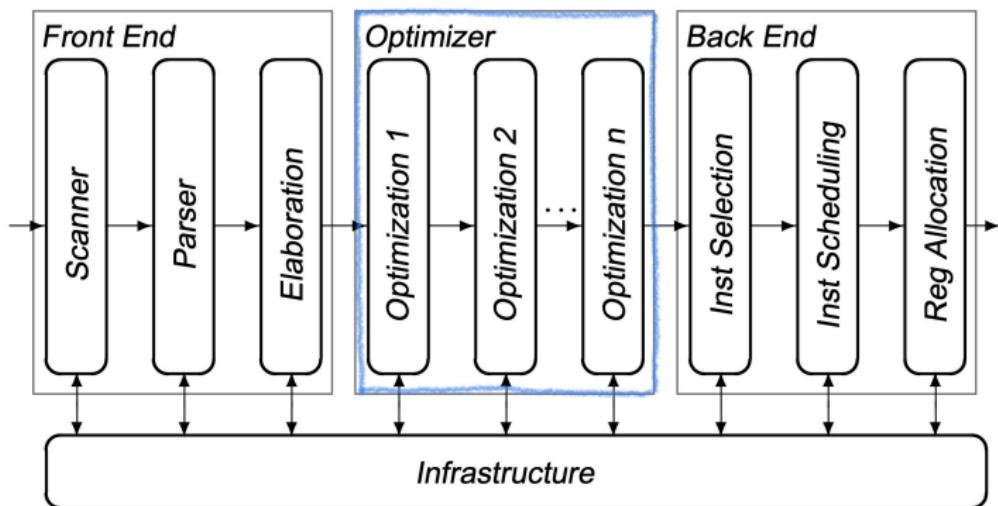
- ▶ In a typical, traditional undergraduate compiler course, each method in your program is compiled **once**.



TIL: Why Is the Interpreter So Interesting?



TIL: Why Is the Interpreter So Interesting?



TIL: Why Is the Interpreter So Interesting?

- ▶ Execution Time = **Jitting Time** + Assembly Execution Time

TIL: Why Is the Interpreter So Interesting?

- ▶ Execution Time = **Jitting Time** + Assembly Execution Time
- ▶ Trade-offs!



TIL: Why Is the Interpreter So Interesting?

- ▶ Execution Time = **Jitting Time** + Assembly Execution Time
- ▶ Trade-offs!



- ▶ **More** optimization means **faster** steady-state performance but **slower** start-up time.

TIL: Why Is the Interpreter So Interesting?

- ▶ Execution Time = **Jitting Time** + Assembly Execution Time
- ▶ Trade-offs!



- ▶ **More** optimization means **faster** steady-state performance but **slower** start-up time.
- ▶ Alternatively, **less** optimization means **faster** start-up time but **slower** steady-state performance.

TIL: Why Is the Interpreter So Interesting?



TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1
- ▶ Multiple compilations for the same method

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1
- ▶ Multiple compilations for the same method
- ▶ Hot-swapped at runtime

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1
- ▶ Multiple compilations for the same method
- ▶ Hot-swapped at runtime
- ▶ **Faster start-up time:** JIT generates the initial compilation quickly, sacrificing code optimization. (**Tier 0, QuickJIT**)

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1
- ▶ Multiple compilations for the same method
- ▶ Hot-swapped at runtime
- ▶ **Faster start-up time:** JIT generates the initial compilation quickly, sacrificing code optimization. (**Tier 0, QuickJIT**)
- ▶ **Faster steady-state performance:** If a method is called frequently, more optimized code is generated **asynchronously**, which then replaces the initial code. (**Tier 1**)

TIL: Why Is the Interpreter So Interesting?

- ▶ **Tiered Compilation!**
- ▶ Introduced in .NET Core 2.1
- ▶ Multiple compilations for the same method
- ▶ Hot-swapped at runtime
- ▶ **Faster start-up time:** JIT generates the initial compilation quickly, sacrificing code optimization. (**Tier 0, QuickJIT**)
- ▶ **Faster steady-state performance:** If a method is called frequently, more optimized code is generated **asynchronously**, which then replaces the initial code. (**Tier 1**)
- ▶ **ReadyToRun** images have several constraints which prohibit some types of optimization. TC generates optimized code that replaces the pre-compiled version.

TIL: Why Is the Interpreter So Interesting?

- ▶ Where do we go from here?

TIL: Why Is the Interpreter So Interesting?

- ▶ Where do we go from here?
- ▶ Execution Time = Jitting Time + Assembly Execution Time

TIL: Why Is the Interpreter So Interesting?

- ▶ Where do we go from here?
- ▶ Execution Time = Jitting Time + Assembly Execution Time
- ▶ Execution Time = Jitting Time + Assembly Execution Time?

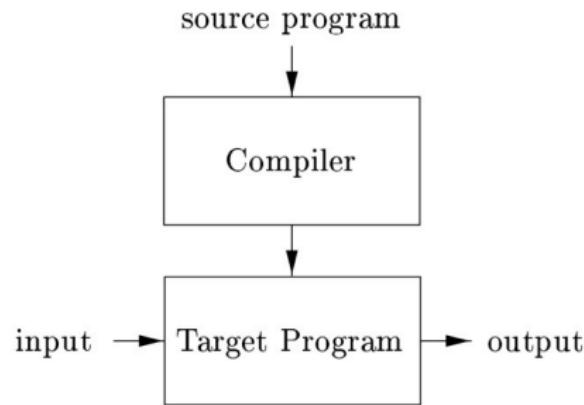
TIL: Why Is the Interpreter So Interesting?

- ▶ Where do we go from here?
- ▶ Execution Time = Jitting Time + Assembly Execution Time
- ▶ Execution Time = Jitting Time + Assembly Execution Time?



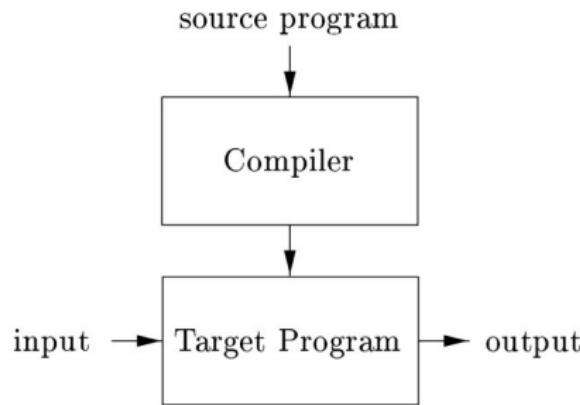
TIL: Why Is the Interpreter So Interesting?

- ▶ This is a **compiler**:

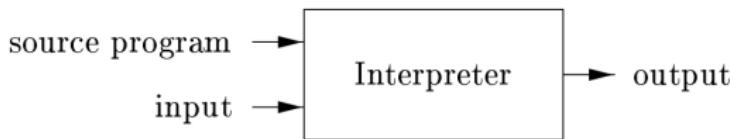


TIL: Why Is the Interpreter So Interesting?

- ▶ This is a **compiler**:



- ▶ This is an **interpreter**:



TIL: Why Is the Interpreter So Interesting?

- ▶ An **interpreter** does not need to JIT assembly code before it is able to execute methods.

TIL: Why Is the Interpreter So Interesting?

- ▶ An **interpreter** does not need to JIT assembly code before it is able to execute methods.
- ▶ Hence, an **interpreter** executes IL code **immediately**.

TIL: Why Is the Interpreter So Interesting?

- ▶ An **interpreter** does not need to JIT assembly code before it is able to execute methods.
- ▶ Hence, an **interpreter** executes IL code **immediately**.
- ▶ Execution Time = Jitting Time + Interpretation Time

TIL: Why Is the Interpreter So Interesting?

- ▶ An **interpreter** does not need to JIT assembly code before it is able to execute methods.
- ▶ Hence, an **interpreter** executes IL code **immediately**.
- ▶ Execution Time = Jitting Time + Interpretation Time

Conjecture

Replacing the current Tier 0 JIT with an interpreter in Tiered Compilation would achieve faster start-up time, while preserving the fast steady-state performance from the Tier 1 JIT, and not deteriorating the total execution time by too much.

Where Are We?

Where Are We?

- ▶ Why Is the Interpreter So Interesting?

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots

Where Are We?

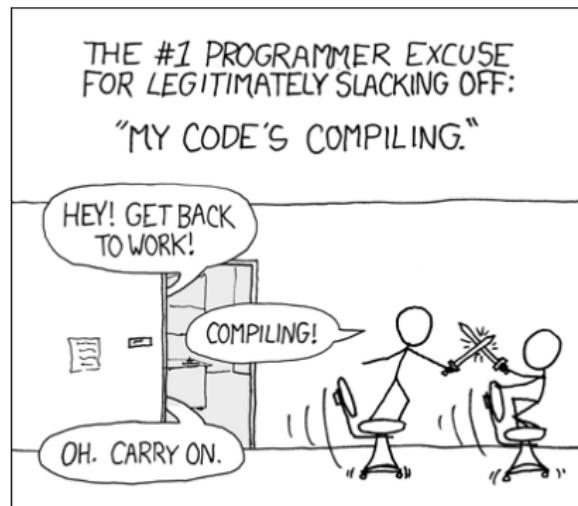
- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks!

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks!
- ▶ Where Do We Go from Here?

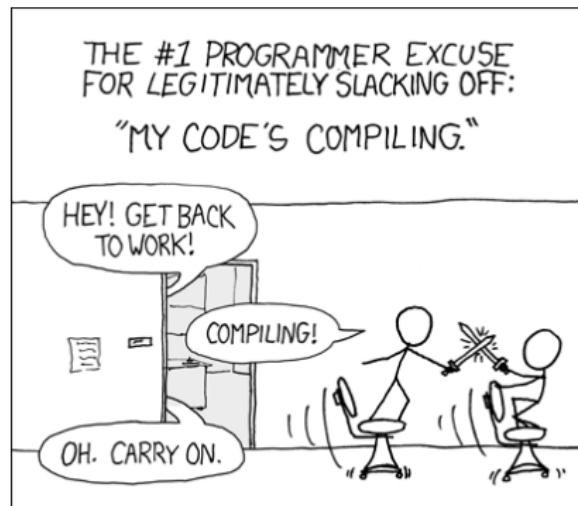
Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks!
- ▶ Where Do We Go from Here?



Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ... ↗
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks!
- ▶ Where Do We Go from Here?



TIL: ... and Fixes, and Fixes, and Fixes, ...

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ The interpreter did not run on HelloWorld in its initial state.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ The interpreter did not run on HelloWorld in its initial state.

```
static bool Not(bool p)
{
    bool q = p;
    q = !q;
    return q;
}

static void Main(string[] args)
{
    Console.WriteLine("Hello World!");
    bool b = true;
    bool bb = Not(b);

    bb = Not(bb);
    // ... 30 additional lines of `bb = Not(bb);` 

    Console.WriteLine(bb);
}
```

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ The interpreter did not run on HelloWorld in its initial state.

```
static bool Not(bool p)
{
    bool q = p;
    q = !q;
    return q;
}

static void Main(string[] args)
{
    Console.WriteLine("Hello World!");
    bool b = true;
    bool bb = Not(b);

    bb = Not(bb);
    // ... 30 additional lines of `bb = Not(bb);` 

    Console.WriteLine(bb);
}
```

- ▶ The number of calls to Not is used to trigger the promotion to Tier 1 in Tiered Compilation.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.
- ▶ **Fix #4:** Provided a workaround for the handling of hardware intrinsics in the interpreter.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.
- ▶ **Fix #4:** Provided a workaround for the handling of hardware intrinsics in the interpreter.
- ▶ **Fix #5:** Two new COM+ environment variables:

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.
- ▶ **Fix #4:** Provided a workaround for the handling of hardware intrinsics in the interpreter.
- ▶ **Fix #5:** Two new COM+ environment variables:
 - ▶ `ForceInterpreter`: Force the interpreter always to be used when executing a program.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.
- ▶ **Fix #4:** Provided a workaround for the handling of hardware intrinsics in the interpreter.
- ▶ **Fix #5:** Two new COM+ environment variables:
 - ▶ ForceInterpreter: Force the interpreter always to be used when executing a program.
 - ▶ InterpreterHWIntrinsicsIsSupportedFalse: Control the workaround in **Fix #4** since it is not production-level.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Fix #1:** Fixed the type cast compilation errors.
- ▶ **Fix #2:** Fixed several bugs in the implementation of the asynchronous tiering-up in Tiered Compilation.
- ▶ **Fix #3:** Fixed the call info of virtual function calls to avoid the “Bad IL Format” exception.
- ▶ **Fix #4:** Provided a workaround for the handling of hardware intrinsics in the interpreter. ←
- ▶ **Fix #5:** Two new COM+ environment variables:
 - ▶ ForceInterpreter: Force the interpreter always to be used when executing a program.
 - ▶ InterpreterHWIntrinsicsIsSupportedFalse: Control the workaround in **Fix #4** since it is not production-level.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Hardware Intrinsics**

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Hardware Intrinsics**
- ▶ **Issue:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Hardware Intrinsics**
- ▶ **Issue:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
    bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call     bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Hardware Intrinsics**
- ▶ **Issue:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
    bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call      bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

- ▶ **“Fix” (a.k.a. Workaround):** Force the interpreter to return `false` for `get_IsSupported()` to disable hardware intrinsics.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Hardware Intrinsics**
- ▶ **Issue:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
    bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call      bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

- ▶ **“Fix” (a.k.a. Workaround):** Force the interpreter to return `false` for `get_IsSupported()` to disable hardware intrinsics.
- ▶ Controlled by the COM+ environment variable
`InterpreterHWIntrinsicsIsSupportedFalse`

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies.
- ▶ **Fix #2:** Implement hardware intrinsics methods in the interpreter.

TIL: ... and Fixes, and Fixes, and Fixes, ...



TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies. ↗
- ▶ **Fix #2:** Implement hardware intrinsics methods in the interpreter.

TIL: ... and Fixes, and Fixes, and Fixes, ...

FIX ALL THE
THINGS!



TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **TIL:** There are too many arbitrary bugs of the interpreter functionalities to be fixed within a 12-week internship.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **TIL:** There are too many arbitrary bugs of the interpreter functionalities to be fixed within a 12-week internship.
- ▶ **TIL:** The performance analysis of the interpreter might be more interesting than fixing functionality bugs.

TIL: ... and Fixes, and Fixes, and Fixes, ...

- ▶ **TIL:** There are too many arbitrary bugs of the interpreter functionalities to be fixed within a 12-week internship.
- ▶ **TIL:** The performance analysis of the interpreter might be more interesting than fixing functionality bugs.
- ▶ **Priority 0 tests:**

Total: 2822

Failed: 184

Passing Percentage: 93.48%

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots ↪
- ▶ Oh My! Bottlenecks!
- ▶ Where Do We Go from Here?



TIL from Beautiful Plots

TIL from Beautiful Plots

Conjecture

Replacing the current Tier 0 JIT with an interpreter in Tiered Compilation would achieve faster start-up time, while preserving the fast steady-state performance from the Tier 1 JIT, and not deteriorating the total execution time by too much.

TIL from Beautiful Plots

Conjecture

Replacing the current Tier 0 JIT with an interpreter in Tiered Compilation would achieve faster start-up time, while preserving the fast steady-state performance from the Tier 1 JIT, and not deteriorating the total execution time by too much.

- ▶ **Steady-state performance:** SciMark 2.0 with BenchmarkDotNet, including benchmarks of scientific and numerical computing, e.g. Fast Fourier Transform (FFT), Successive Over-Relaxation (SOR), Monte Carlo Integration, Sparse Matrix Multiply, and LU Matrix Factorization.

TIL from Beautiful Plots

Conjecture

Replacing the current Tier 0 JIT with an interpreter in Tiered Compilation would achieve faster start-up time, while preserving the fast steady-state performance from the Tier 1 JIT, and not deteriorating the total execution time by too much.

- ▶ **Steady-state performance:** SciMark 2.0 with BenchmarkDotNet, including benchmarks of scientific and numerical computing, e.g. Fast Fourier Transform (FFT), Successive Over-Relaxation (SOR), Monte Carlo Integration, Sparse Matrix Multiply, and LU Matrix Factorization.
- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?

TIL from Beautiful Plots

Conjecture

Replacing the current Tier 0 JIT with an interpreter in Tiered Compilation would achieve faster start-up time, while preserving the fast steady-state performance from the Tier 1 JIT, and not deteriorating the total execution time by too much.

- ▶ **Steady-state performance:** SciMark 2.0 with BenchmarkDotNet, including benchmarks of scientific and numerical computing, e.g. Fast Fourier Transform (FFT), Successive Over-Relaxation (SOR), Monte Carlo Integration, Sparse Matrix Multiply, and LU Matrix Factorization.
- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?
- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?

TIL from Beautiful Plots

- ▶ Five configurations:

TIL from Beautiful Plots

- ▶ Five configurations:
- ▶ default: the current default configuration of Tiered Compilation, with a tiering-up threshold of 30

TIL from Beautiful Plots

- ▶ Five configurations:
- ▶ default: the current default configuration of Tiered Compilation, with a tiering-up threshold of 30
- ▶ jit-min-opt: always JIT and execute the minimally optimized assembly code (Tier 0)

TIL from Beautiful Plots

- ▶ Five configurations:
- ▶ default: the current default configuration of Tiered Compilation, with a tiering-up threshold of 30
- ▶ jit-min-opt: always JIT and execute the minimally optimized assembly code (Tier 0)
- ▶ tier1-only: always JIT and execute the highly optimized assembly code (Tier 1)

TIL from Beautiful Plots

- ▶ Five configurations:
- ▶ default: the current default configuration of Tiered Compilation, with a tiering-up threshold of 30
- ▶ jit-min-opt: always JIT and execute the minimally optimized assembly code (Tier 0)
- ▶ tier1-only: always JIT and execute the highly optimized assembly code (Tier 1)
- ▶ interp-only: always interpret IL code with the interpreter

TIL from Beautiful Plots

- ▶ Five configurations:
- ▶ default: the current default configuration of Tiered Compilation, with a tiering-up threshold of 30
- ▶ jit-min-opt: always JIT and execute the minimally optimized assembly code (Tier 0)
- ▶ tier1-only: always JIT and execute the highly optimized assembly code (Tier 1)
- ▶ interp-only: always interpret IL code with the interpreter
- ▶ interp-tc: Tiered Compilation with Tier 0 substituted with the interpreter, with a tiering-up threshold of 0

TIL from Beautiful Plots

- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?

TIL from Beautiful Plots

- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?
- ▶ Computed the average execution time over 15-20 iterations for each benchmark, after throwing away warm-up iterations

TIL from Beautiful Plots

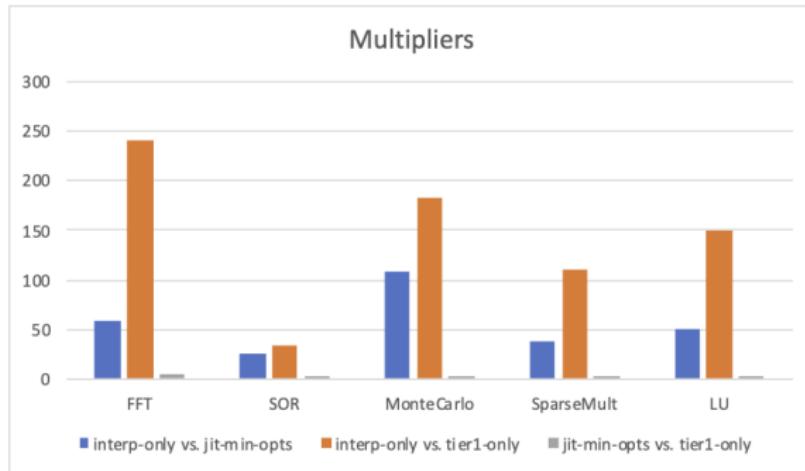
- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?
- ▶ Computed the average execution time over 15-20 iterations for each benchmark, after throwing away warm-up iterations

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844

TIL from Beautiful Plots

- ▶ **Question #1:** How much is the interpreter slower than the Tier 0 and Tier 1 JIT, respectively?
- ▶ Computed the average execution time over 15-20 iterations for each benchmark, after throwing away warm-up iterations

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844



TIL from Beautiful Plots

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844

TIL from Beautiful Plots

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844

- ▶ Compared to QuickJIT, the interpreter is 25-times slower in the best case (SOR), and 108-times slower in the worst case (MonteCarlo).

TIL from Beautiful Plots

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844

- ▶ Compared to QuickJIT, the interpreter is 25-times slower in the best case (SOR), and 108-times slower in the worst case (MonteCarlo).
- ▶ Compared to Tier 1, the interpreter is 34-times slower in the best case (SOR), and 241-times slower in the worst case (FFT).

TIL from Beautiful Plots

	interp-only vs. jit-min-opts	interp-only vs. tier1-only	jit-min-opts vs. tier1-only
FFT	58.51598622	241.0419967	4.119250349
SOR	25.35198702	33.61917818	1.326096379
MonteCarlo	107.8597921	182.8229512	1.695005596
SparseMult	37.65475619	109.7832527	2.915521539
LU	49.93685214	148.7842465	2.979447844

- ▶ Compared to QuickJIT, the interpreter is 25-times slower in the best case (SOR), and 108-times slower in the worst case (MonteCarlo).
- ▶ Compared to Tier 1, the interpreter is 34-times slower in the best case (SOR), and 241-times slower in the worst case (FFT).

TIL from Beautiful Plots

- ▶ **Accidentally**, also measured the steady-state performance for each benchmark with a **Checked** build.

TIL from Beautiful Plots

- ▶ **Accidentally**, also measured the steady-state performance for each benchmark with a **Checked** build.

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	473.5579762	1806.993473	3.815780884
SOR	333.4660434	440.5411904	1.321097602
MonteCarlc	2306.499301	3937.286102	1.707039798
SparseMult	410.4487699	1213.298956	2.956030191
LU	380.695905	1595.84787	4.191922868

TIL from Beautiful Plots

- ▶ **Accidentally**, also measured the steady-state performance for each benchmark with a **Checked** build.

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	473.5579762	1806.993473	3.815780884
SOR	333.4660434	440.5411904	1.321097602
MonteCarlo	2306.499301	3937.286102	1.707039798
SparseMult	410.4487699	1213.298956	2.956030191
LU	380.695905	1595.84787	4.191922868

Observation

The multipliers with a **Checked** build is between 333–2306 compared to `jit-min-opt`, and between 441–3937 compared to `tier1-only`, much greater than a **Release** build.

TIL from Beautiful Plots

- ▶ **Accidentally**, also measured the steady-state performance for each benchmark with a **Checked** build.

	interp-only vs. jit-min-opt	interp-only vs. tier1-only	jit-min-opt vs. tier1-only
FFT	473.5579762	1806.993473	3.815780884
SOR	333.4660434	440.5411904	1.321097602
MonteCarlo	2306.499301	3937.286102	1.707039798
SparseMult	410.4487699	1213.298956	2.956030191
LU	380.695905	1595.84787	4.191922868

Observation

The multipliers with a **Checked** build is between 333–2306 compared to `jit-min-opt`, and between 441–3937 compared to `tier1-only`, much greater than a **Release** build.

Conjecture

The interpretation of IL code is the execution of native C++ code. A **Release** build disables many assertions and enables many optimizations of the native code.

TIL from Beautiful Plots

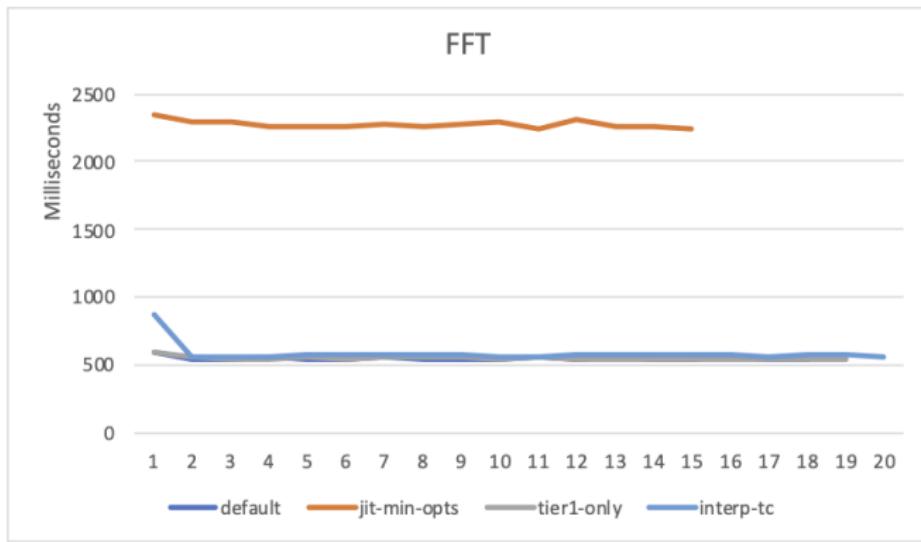
- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?

TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.

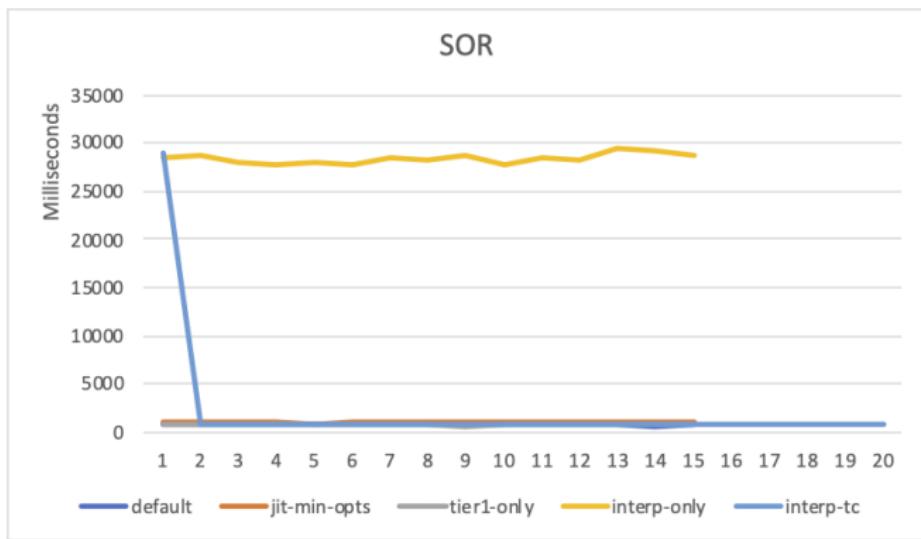
TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.



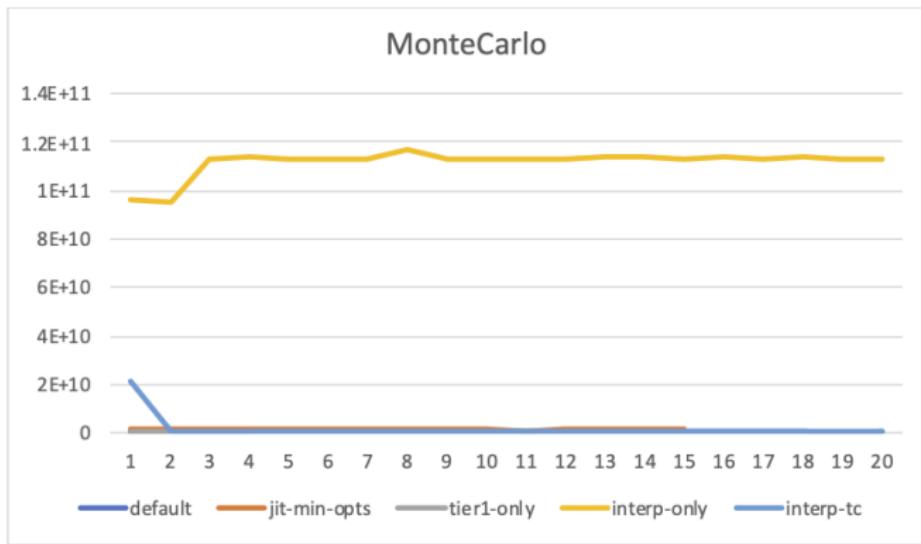
TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.



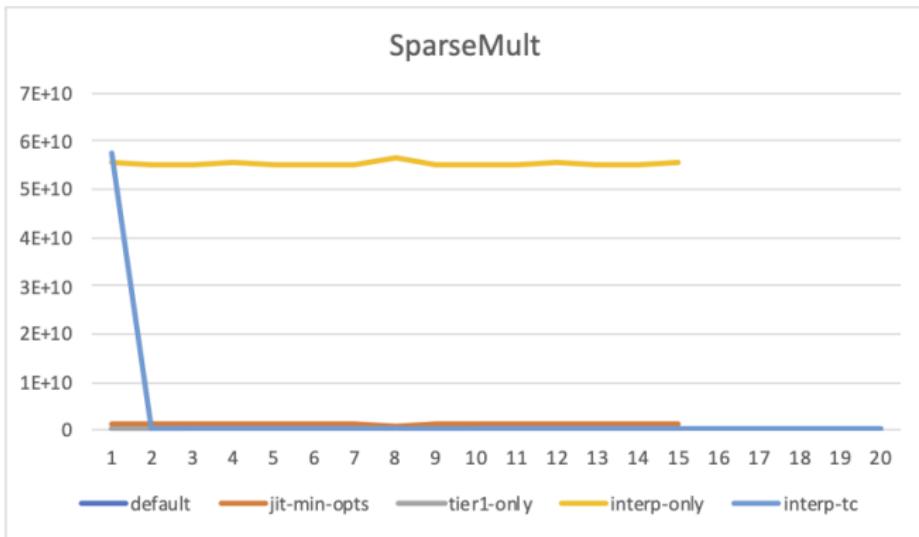
TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.



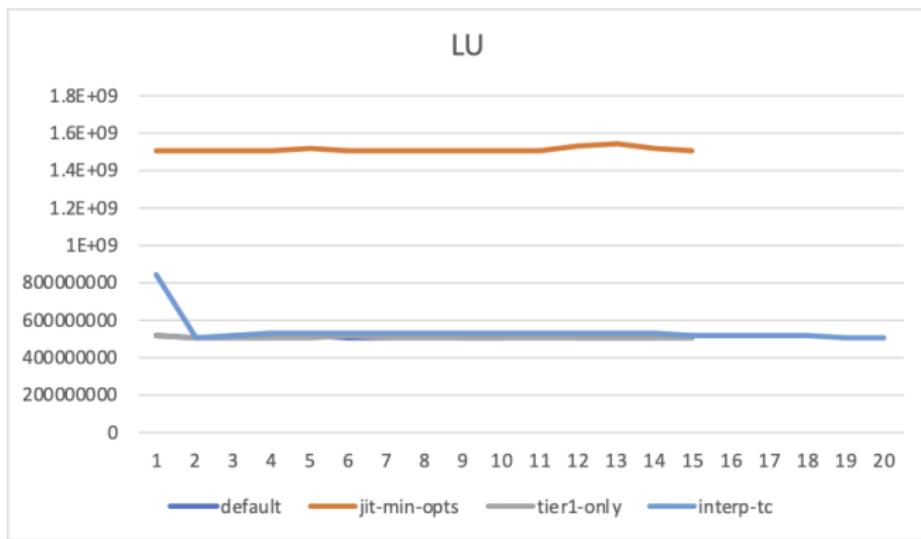
TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.



TIL from Beautiful Plots

- ▶ **Question #2:** Is the steady-state performance after several warm-up iterations close to the Tier 1 JIT performance?
- ▶ Yes.



TIL from Beautiful Plots

- ▶ Start-up time

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence $1, 2, 4, 8, 16, \dots$. Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence $1, 2, 4, 8, 16, \dots$. Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)
- ▶ Two additional configurations:

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence 1, 2, 4, 8, 16, Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)
- ▶ Two additional configurations:
- ▶ `interp-without-loops`: interpret IL code with the interpreter, except that methods containing loops are JIT'ed

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence 1, 2, 4, 8, 16, Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)
- ▶ Two additional configurations:
- ▶ `interp-without-loops`: interpret IL code with the interpreter, except that methods containing loops are JIT'ed
- ▶ `interp-tc-without-loops`: Tiered Compilation with Tier 0 substituted with the interpreter with a tiering-up threshold of 0, except that methods containing loops are JIT'ed

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence 1, 2, 4, 8, 16, Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)
- ▶ Two additional configurations:
- ▶ `interp-without-loops`: interpret IL code with the interpreter, except that methods containing loops are JIT'ed
- ▶ `interp-tc-without-loops`: Tiered Compilation with Tier 0 substituted with the interpreter with a tiering-up threshold of 0, except that methods containing loops are JIT'ed
- ▶ **Question #1:** Does the interpreter easily beat QuickJIT and Tier 1 JIT for small iteration counts?

TIL from Beautiful Plots

- ▶ **Start-up time**
- ▶ A computational loop, where the number of iterations taken from an exponential sequence 1, 2, 4, 8, 16, Plot the total execution time of the computational loop against the number of iterations. (@BruceForstall @AndyAyersMS)
- ▶ Two additional configurations:
- ▶ `interp-without-loops`: interpret IL code with the interpreter, except that methods containing loops are JIT'ed
- ▶ `interp-tc-without-loops`: Tiered Compilation with Tier 0 substituted with the interpreter with a tiering-up threshold of 0, except that methods containing loops are JIT'ed
- ▶ **Question #1:** Does the interpreter easily beat QuickJIT and Tier 1 JIT for small iteration counts?
- ▶ **Question #2:** What are the **crossover** times when QuickJIT beats the interpreter, and when Tier 1 JIT beats QuickJIT?

TIL from Beautiful Plots

```
[MethodImplAttribute(MethodImplOptions.NoInlining)]
private static void update(ref int a, ref int b)
{
    int c = a + b;

    if (c >= MOD)
    {
        c -= MOD;
    }

    a = b;
    b = c;
}

public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}
```

TIL from Beautiful Plots

```
[MethodImplAttribute(MethodImplOptions.NoInlining)]
public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}

private static void update(ref int a, ref int b)
{
    int c = a + b;

    if (c >= MOD)
    {
        c -= MOD;
    }

    a = b;
    b = c;
}
```

- ▶ **execute** is only invoked **once** for each benchmark, and hence will not be tiered up.

TIL from Beautiful Plots

```
[MethodImplAttribute(MethodImplOptions.NoInlining)]
public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}

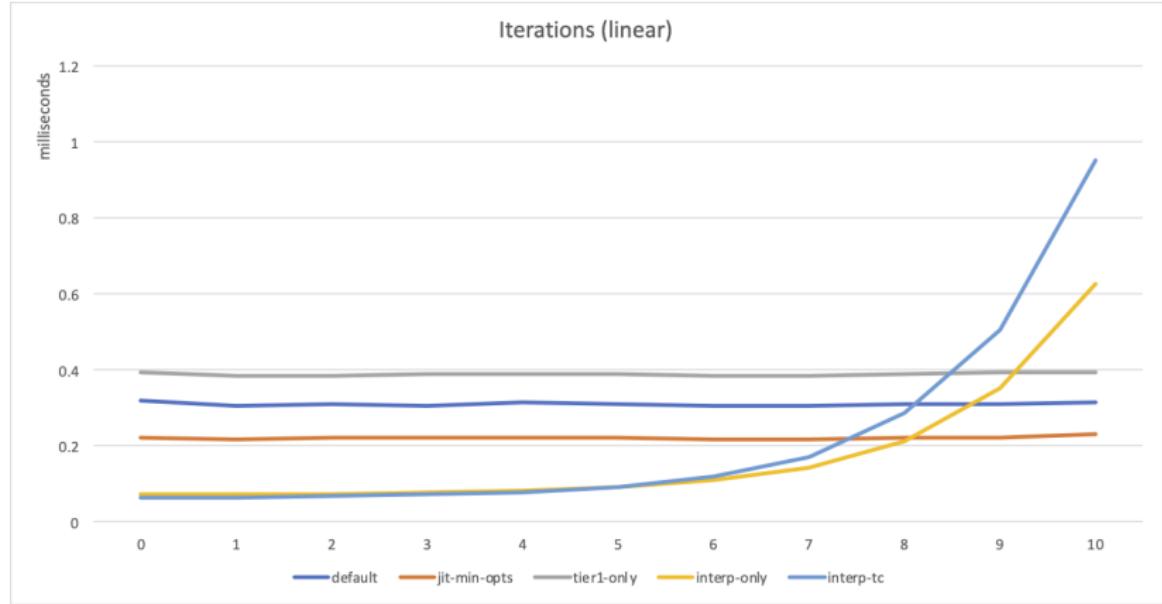
private static void update(ref int a, ref int b)
{
    int c = a + b;

    if (c >= MOD)
    {
        c -= MOD;
    }

    a = b;
    b = c;
}
```

- ▶ `execute` is only invoked **once** for each benchmark, and hence will not be tiered up.
- ▶ Set `COMPlus_InterpreterDoLoopMethods=0` to JIT execute while interpreting `update`.

TIL from Beautiful Plots



TIL from Beautiful Plots

- ▶ The JIT curves are **flat** for small iteration counts. The *y*-values of the flat curves correspond to the jitting time of different configurations.

TIL from Beautiful Plots

- ▶ The JIT curves are **flat** for small iteration counts. The y -values of the flat curves correspond to the jitting time of different configurations.
- ▶ The execution time of a single iteration approximately equals the start-up time of the interpreter.

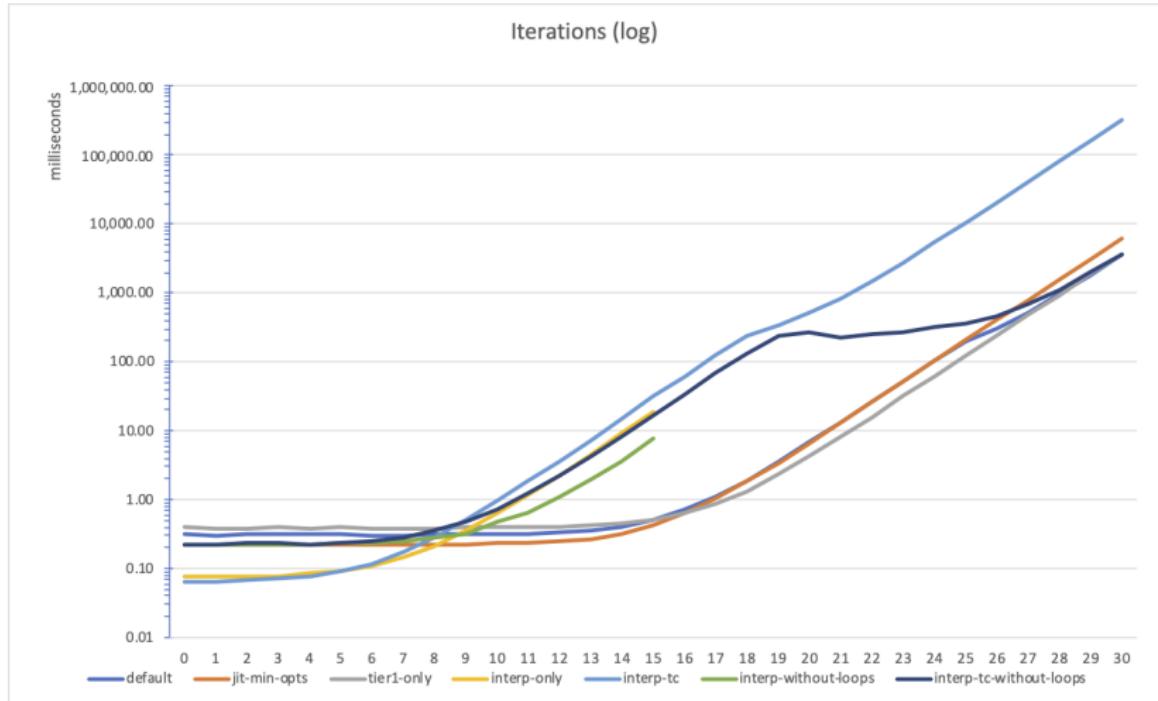
TIL from Beautiful Plots

- ▶ The JIT curves are **flat** for small iteration counts. The y -values of the flat curves correspond to the jitting time of different configurations.
- ▶ The execution time of a single iteration approximately equals the start-up time of the interpreter.
- ▶ **TIL:** The start-up performance of the interpreter (≈ 0.075 ms) is approximately 3 times that of QuickJIT (≈ 0.22 ms), and 5 times that of Tier 1 JIT (≈ 0.39 ms).

TIL from Beautiful Plots

- ▶ The JIT curves are **flat** for small iteration counts. The y -values of the flat curves correspond to the jitting time of different configurations.
- ▶ The execution time of a single iteration approximately equals the start-up time of the interpreter.
- ▶ **TIL:** The start-up performance of the interpreter (≈ 0.075 ms) is approximately 3 times that of QuickJIT (≈ 0.22 ms), and 5 times that of Tier 1 JIT (≈ 0.39 ms).
- ▶ **TIL:** The interpreter beats QuickJIT for the first 2^8 iterations, and beats Tier 1 JIT for the first 2^9 iteration.

TIL from Beautiful Plots



TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter?

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter? What about a combination of synchronous and asynchronous tiering-up?

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter? What about a combination of synchronous and asynchronous tiering-up?
- ▶ **Question:** Background requests go into a queue, building up a big backlog. Can this be turned into a priority queue? What criteria to determine the priority?

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter? What about a combination of synchronous and asynchronous tiering-up?
- ▶ **Question:** Background requests go into a queue, building up a big backlog. Can this be turned into a priority queue? What criteria to determine the priority?
- ▶ A “turning point” at the 2^{18} th iteration on interp-tc.

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter? What about a combination of synchronous and asynchronous tiering-up?
- ▶ **Question:** Background requests go into a queue, building up a big backlog. Can this be turned into a priority queue? What criteria to determine the priority?
- ▶ A “turning point” at the 2^{18} th iteration on interp-tc.
- ▶ **TIL:** update gets tiered up. However, the total execution time is dominated by the interpretation time of the loop, which cannot be tiered up.

TIL from Beautiful Plots

- ▶ default gets tiered up at the 2^{25} th iteration. (threshold: 30)
- ▶ interp-tc and interp-tc-without-loops get tiered up at the 2^{18} th and 2^{19} th iterations, respectively. (thresholds: 0)
- ▶ **TIL:** The **actual, asynchronous** tiering-up occurs much later after it is queued in the background.
- ▶ **Question:** Would **synchronous** tiering-up have better overall performance if Tier 0 is the interpreter? What about a combination of synchronous and asynchronous tiering-up?
- ▶ **Question:** Background requests go into a queue, building up a big backlog. Can this be turned into a priority queue? What criteria to determine the priority?
- ▶ A “turning point” at the 2^{18} th iteration on interp-tc.
- ▶ **TIL:** update gets tiered up. However, the total execution time is dominated by the interpretation time of the loop, which cannot be tiered up. ←

TIL from Beautiful Plots



TIL: On Stack Replacement

TIL: On Stack Replacement

- ▶ **Tiered Compilation:** If a method is invoked frequently, optimize it.

TIL: On Stack Replacement

- ▶ **Tiered Compilation:** If a method is invoked frequently, optimize it.

```
public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}
```

TIL: On Stack Replacement

- ▶ **Tiered Compilation:** If a method is invoked frequently, optimize it.

```
public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}
```

- ▶ If a method is invoked very few times but its execution takes long time, we can never optimize it with Tiered Compilation **only**.

TIL: On Stack Replacement

- ▶ **Tiered Compilation:** If a method is invoked frequently, optimize it.

```
public static int execute(int num_iterations)
{
    int i;
    int a = 0, b = 1;

    for (i = 0; i < num_iterations; i++)
    {
        update(ref a, ref b);
    }

    return a;
}
```

- ▶ If a method is invoked very few times but its execution takes long time, we can never optimize it with Tiered Compilation **only**.
- ▶ COMPlus_InterpreterDoLoopMethods=0 to avoid interpreting methods with loops.

TIL: On Stack Replacement

- ▶ **On Stack Replacement (OSR):** If some part of a method is called frequently, optimize that part.

TIL: On Stack Replacement

- ▶ **On Stack Replacement (OSR):** If some part of a method is called frequently, optimize that part.
- ▶ Formally, OSR allows the code executed by currently running methods to be changed in the middle of method execution, while those methods are active "on stack."

TIL: On Stack Replacement

- ▶ **On Stack Replacement (OSR):** If some part of a method is called frequently, optimize that part.
- ▶ Formally, OSR allows the code executed by currently running methods to be changed in the middle of method execution, while those methods are active "on stack."
- ▶ With OSR, we can safely start with interpreting a method with a loop. If the loop body is executed frequently, OSR allows the loop to be optimized.

TIL: On Stack Replacement

- ▶ **On Stack Replacement (OSR):** If some part of a method is called frequently, optimize that part.
- ▶ Formally, OSR allows the code executed by currently running methods to be changed in the middle of method execution, while those methods are active "on stack."
- ▶ With OSR, we can safely start with interpreting a method with a loop. If the loop body is executed frequently, OSR allows the loop to be optimized.
- ▶ A **patchpoint** is a particular location in a method that supports OSR transitions.

TIL: On Stack Replacement

- ▶ **On Stack Replacement (OSR):** If some part of a method is called frequently, optimize that part.
- ▶ Formally, OSR allows the code executed by currently running methods to be changed in the middle of method execution, while those methods are active "on stack."
- ▶ With OSR, we can safely start with interpreting a method with a loop. If the loop body is executed frequently, OSR allows the loop to be optimized.
- ▶ A **patchpoint** is a particular location in a method that supports OSR transitions.
- ▶ At a patchpoint, the **live state** of the ongoing computation must be identifiable (e.g. registers, stack slots, return address, implicit arguments).

TIL: On Stack Replacement

- ▶ Commonly, targets of **lexical** loop back edges or backward branches are chosen to be patchpoints.

TIL: On Stack Replacement

- ▶ Commonly, targets of **lexical** loop back edges or backward branches are chosen to be patchpoints.
- ▶ Identified by a single scan of the IL code of a method

TIL: On Stack Replacement

- ▶ Commonly, targets of **lexical** loop back edges or backward branches are chosen to be patchpoints.
- ▶ Identified by a single scan of the IL code of a method
- ▶ **Partially interruptible**

TIL: On Stack Replacement

- ▶ Commonly, targets of **lexical** loop back edges or backward branches are chosen to be patchpoints.
- ▶ Identified by a single scan of the IL code of a method
- ▶ **Partially interruptible**
- ▶ With such choices, a method cannot execute indefinitely between patchpoints.

TIL: On Stack Replacement

- ▶ Commonly, targets of **lexical** loop back edges or backward branches are chosen to be patchpoints.
- ▶ Identified by a single scan of the IL code of a method
- ▶ **Partially interruptible**
- ▶ With such choices, a method cannot execute indefinitely between patchpoints.
- ▶ Loops that execute with **non-empty evaluation stacks** are likely rare, hence minimizing and regularizing the live state.

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state
- ▶ **Option #2:** Copy the live state and create a new patchpoint information structure

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state
- ▶ **Option #2:** Copy the live state and create a new patchpoint information structure
- ▶ QuickJIT currently uses **Option #1**.

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state
- ▶ **Option #2:** Copy the live state and create a new patchpoint information structure
- ▶ QuickJIT currently uses **Option #1**.
- ▶ In the interpreter, the local variables and parameters are stored in a custom structure.

TIL: On Stack Replacement

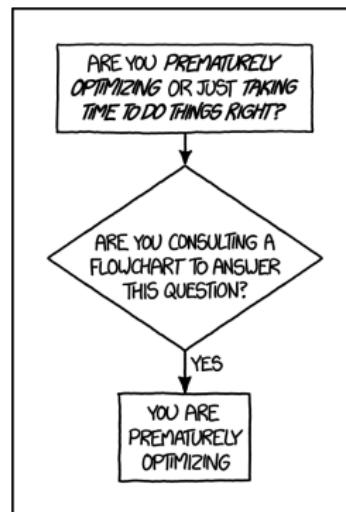
- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state
- ▶ **Option #2:** Copy the live state and create a new patchpoint information structure
- ▶ QuickJIT currently uses **Option #1**.
- ▶ In the interpreter, the local variables and parameters are stored in a custom structure.
- ▶ Optimized jitted code needs to access these pieces of memory.

TIL: On Stack Replacement

- ▶ How does OSR work if the interpreter is used as Tier 0?
- ▶ **Challenge:** The current approach of capturing the live state is limited to JIT'ed code. Local variables and parameters are stored in a custom structure in the interpreter.
- ▶ **Option #1:** Pointers to the piece of memory that contains the live state
- ▶ **Option #2:** Copy the live state and create a new patchpoint information structure
- ▶ QuickJIT currently uses **Option #1**.
- ▶ In the interpreter, the local variables and parameters are stored in a custom structure.
- ▶ Optimized jitted code needs to access these pieces of memory.
- ▶ **Option #2** might be easier to implement without updates to the original pieces of memory.

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks! ←
- ▶ Where Do We Go from Here?



TIL: Oh My! Bottlenecks!

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis
- ▶ **inclusive time:** the execution time including its callees

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis
- ▶ **inclusive time:** the execution time including its callees
- ▶ **exclusive time:** the execution time excluding its callees

TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis
- ▶ **inclusive time:** the execution time including its callees
- ▶ **exclusive time:** the execution time excluding its callees
- ▶ Two approaches:

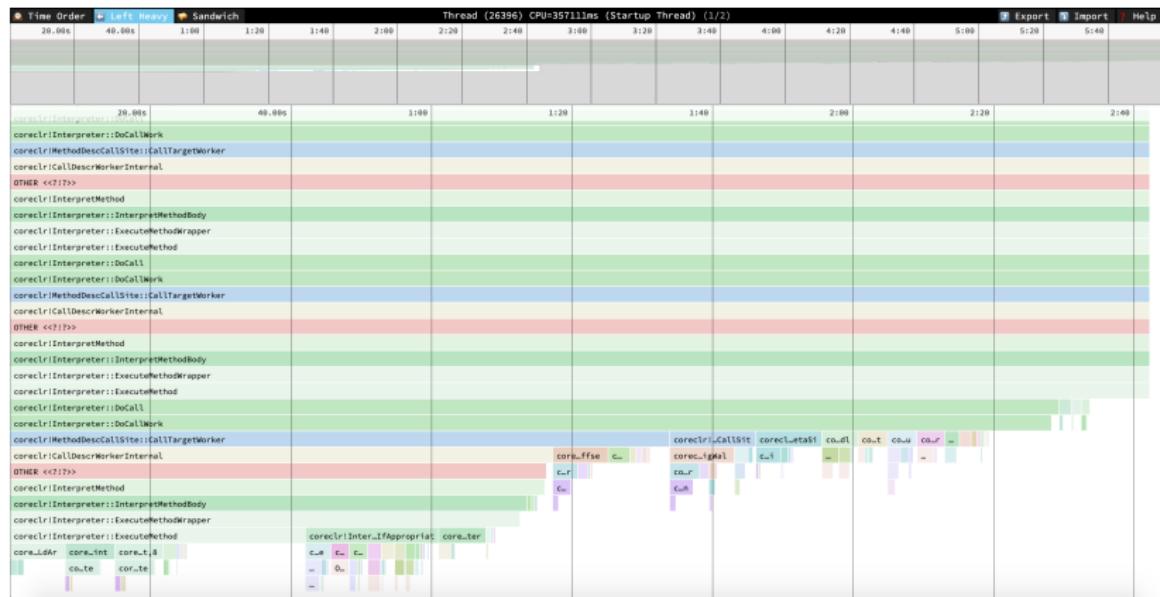
TIL: Oh My! Bottlenecks!

- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis
- ▶ **inclusive time:** the execution time including its callees
- ▶ **exclusive time:** the execution time excluding its callees
- ▶ Two approaches:
- ▶ **Bottom-up:** Start with functions of high exclusive time, and analyze its callers.

TIL: Oh My! Bottlenecks!

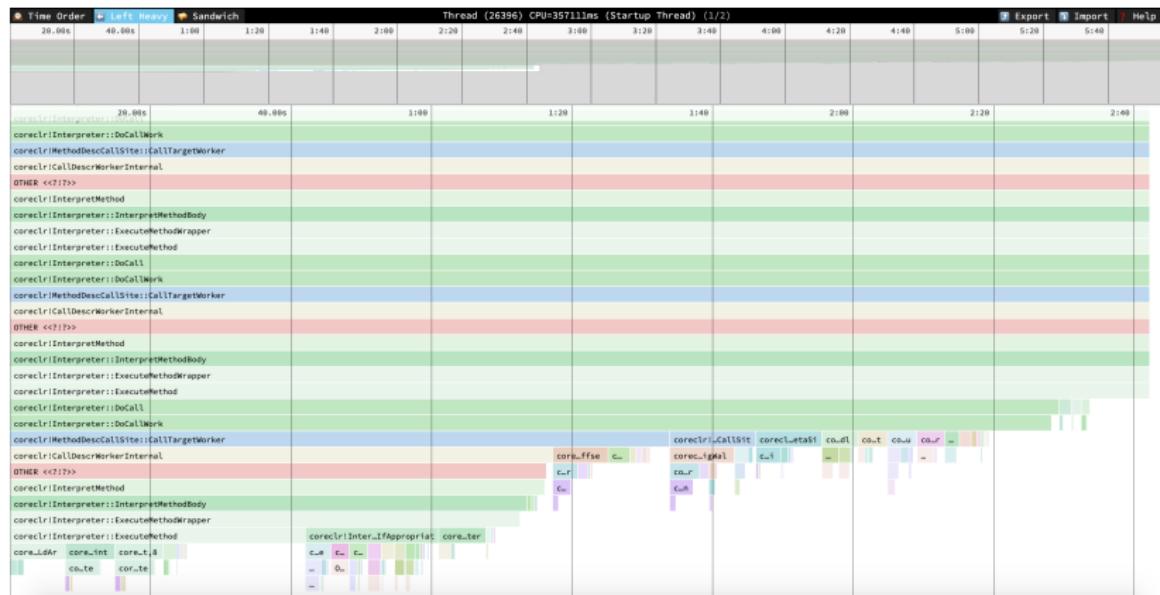
- ▶ **Question:** What makes the interpreter slow?
- ▶ Performance analysis on the ETW profiling data with PerfView and speedscope
- ▶ A very preliminary analysis
- ▶ **inclusive time:** the execution time including its callees
- ▶ **exclusive time:** the execution time excluding its callees
- ▶ Two approaches:
- ▶ **Bottom-up:** Start with functions of high exclusive time, and analyze its callers.
- ▶ **Top-down:** Start from functions of high inclusive time, and analyze its callees to see where the interpreter spends time

TIL: Oh My! Bottlenecks!



- ▶ The pattern of the call stack suggests that DoCall and DoCallWork take a large portion of inclusive time.

TIL: Oh My! Bottlenecks!

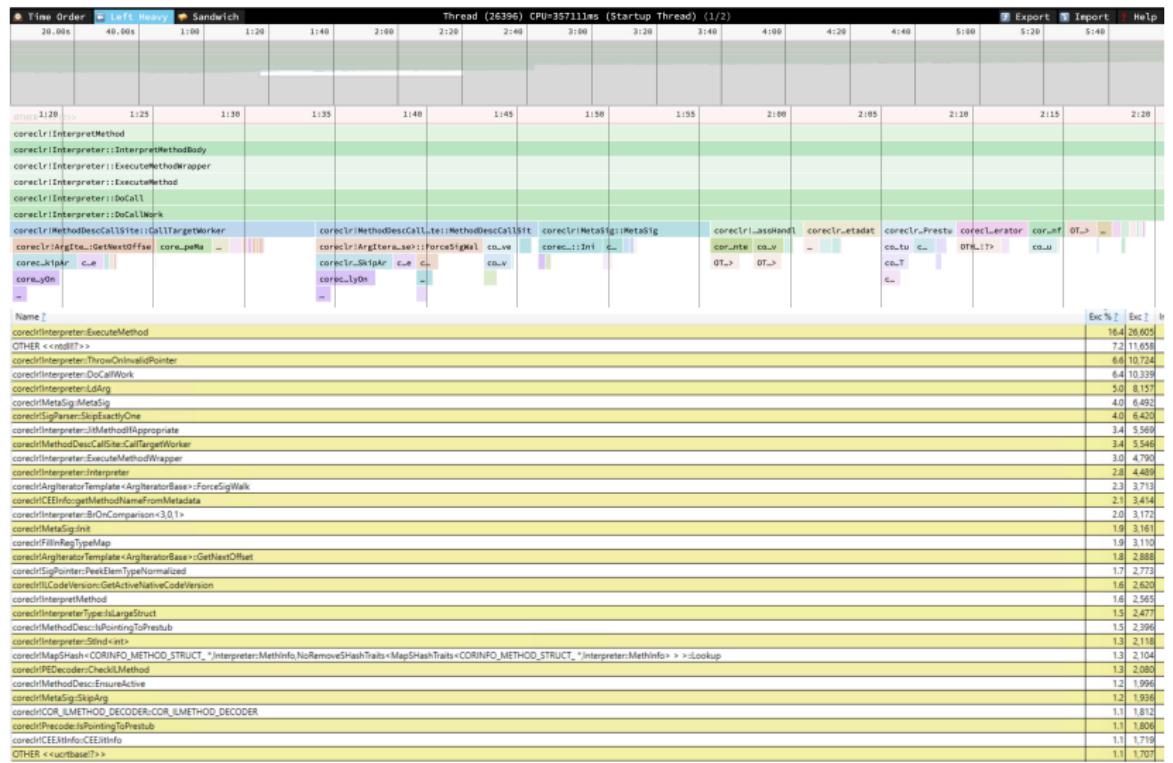


- ▶ The pattern of the call stack suggests that DoCall and DoCallWork take a large portion of inclusive time.
- ▶ **Question:** Which callees (direct or indirect) of DoCallWork have large inclusive and exclusive time?

TIL: Oh My! Bottlenecks!



TIL: Oh My! Bottlenecks!



TIL: Oh My! Bottlenecks!

- ▶ **TIL:** The flamegraph and the exclusive time data suggest that the interpretation of function calls spends much time on parsing the signature and the method descriptor of the callee.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** The flamegraph and the exclusive time data suggest that the interpretation of function calls spends much time on parsing the signature and the method descriptor of the callee.
- ▶ e.g. `MetaSig::MegaSig` (3.7%),
`SigParse::SkipExactlyOnce` (3.6%),
`ArgIteratorTemplate::ForceSigWalk` (2.2%)
`MegaSig::Init` (1.8%), `FillInRegTypeMap` (1.8%) etc.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** The flamegraph and the exclusive time data suggest that the interpretation of function calls spends much time on parsing the signature and the method descriptor of the callee.
- ▶ e.g. `MetaSig::MegaSig` (3.7%),
`SigParse::SkipExactlyOnce` (3.6%),
`ArgIteratorTemplate::ForceSigWalk` (2.2%)
`MegaSig::Init` (1.8%), `FillInRegTypeMap` (1.8%) etc.
- ▶ The source code of `DoCallWork` also suggests the parsing of the callee signature.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** The flamegraph and the exclusive time data suggest that the interpretation of function calls spends much time on parsing the signature and the method descriptor of the callee.
- ▶ e.g. `MetaSig::MegaSig` (3.7%),
`SigParse::SkipExactlyOnce` (3.6%),
`ArgIteratorTemplate::ForceSigWalk` (2.2%)
`MegaSig::Init` (1.8%), `FillInRegTypeMap` (1.8%) etc.
- ▶ The source code of `DoCallWork` also suggests the parsing of the callee signature.
- ▶ `ForceSigWalk` handles calling conventions for callees, and enumerates each argument of a signature in turn.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** The flamegraph and the exclusive time data suggest that the interpretation of function calls spends much time on parsing the signature and the method descriptor of the callee.
- ▶ e.g. `MetaSig::MegaSig` (3.7%),
`SigParse::SkipExactlyOnce` (3.6%),
`ArgIteratorTemplate::ForceSigWalk` (2.2%)
`MegaSig::Init` (1.8%), `FillInRegTypeMap` (1.8%) etc.
- ▶ The source code of `DoCallWork` also suggests the parsing of the callee signature.
- ▶ `ForceSigWalk` handles calling conventions for callees, and enumerates each argument of a signature in turn.
- ▶ **Question:** Can this information of mapping arguments to registers be cached somehow? If we repeatedly call a method, can we let the arguments stay in the registers?

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** DoCallWork also has large exclusive time of 6.4%.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** DoCallWork also has large exclusive time of 6.4%.
- ▶ An inspection into the implementation of DoCallWork suggests that it re-processes arguments every time a function is called.

TIL: Oh My! Bottlenecks!

- ▶ **TIL:** DoCallWork also has large exclusive time of 6.4%.
- ▶ An inspection into the implementation of DoCallWork suggests that it re-processes arguments every time a function is called.
- ▶ **Question:** Can this information also be cached somehow?

TIL: Oh My! Bottlenecks!

```
// CYCLE PROFILE: BEFORE ARG PROCESSING.

if (sigInfo.hasThis())
{
    if (_callThisArg != NULL)
    {
        if (size_t(_callThisArg) == 0x1)
        {
            args[curArgSlot] = NULL;
        }
        else
        {
            args[curArgSlot] = PtrToArgSlot(_callThisArg);
        }
        argTypes[curArgSlot] = InterpreterType(CORINFO_TYPE_BYREF);
    }
    else
    {
        args[curArgSlot] = PtrToArgSlot(OpStackGet<void*>(argsBase + arg));
        argTypes[curArgSlot] = OpStackTypeGet(argsBase + arg);
        arg++;
    }
    // AV -> NullRef translation is NYI for the interpreter,
    // so we should manually check and throw the correct exception.
    if (args[curArgSlot] == NULL)
    {
        // If we're calling a constructor, we bypass this check since the runtime
        // should have thrown OOM if it was unable to allocate an instance.
        if (_callThisArg == NULL)
        {
            _ASSERT(!methToCall->IsStatic());
            ThrowNullPointerException();
        }
        // ...except in the case of strings, which are both
        // allocated and initialized by their special constructor.
        else
        {
            _ASSERT(methToCall->IsCtor() && methToCall->GetMethodTable()->IsString());
        }
    }
    curArgSlot++;
}
```

TIL: Oh My! Bottlenecks!

- ▶ We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

TIL: Oh My! Bottlenecks!

- We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

```
// Hardware intrinsics are recognized by name.  
const char* namespaceName = NULL;  
const char* className = NULL;  
const char* methodName = m_interpCeeInfo.getMethodNameFromMetadata((CORINFO_METHOD_HANDLE)methToCall, &className, &namespaceName, NULL);  
if (  
#if defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.X86") == 0 &&  
#elif defined(TARGET_ARM64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.Arm") == 0 &&  
#endif // defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(methodName, "get_IsSupported") == 0  
)  
{  
    GCX_COOP();  
    DoGetIsSupported();  
    didIntrinsic = true;  
}
```

TIL: Oh My! Bottlenecks!

- We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

```
// Hardware intrinsics are recognized by name.  
const char* namespaceName = NULL;  
const char* className = NULL;  
const char* methodName = m_interpCeeInfo.getMethodNameFromMetadata((CORINFO_METHOD_HANDLE)methToCall, &className, &namespaceName, NULL);  
if (  
#if defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.X86") == 0 &&  
#elif defined(TARGET_ARM64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.Arm") == 0 &&  
#endif // defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(methodName, "get_IsSupported") == 0  
)  
{  
    GCX_COOP();  
    DoGetIsSupported();  
    didIntrinsic = true;  
}
```

- TIL: Obtaining names from metadata is **expensive**.

TIL: Oh My! Bottlenecks!

- ▶ We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

```
// Hardware intrinsics are recognized by name.  
const char* namespaceName = NULL;  
const char* className = NULL;  
const char* methodName = m_interpCeeInfo.getMethodNameFromMetadata((CORINFO_METHOD_HANDLE)methToCall, &className, &namespaceName, NULL);  
if (  
#if defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.X86") == 0 &&  
#elif defined(TARGET_ARM64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.Arm") == 0 &&  
#endif // defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(methodName, "get_IsSupported") == 0  
)  
{  
    GCX_COOP();  
    DoGetIsSupported();  
    didIntrinsic = true;  
}
```

- ▶ **TIL:** Obtaining names from metadata is **expensive**.
- ▶ **CEEInfo::getMethodNameFromMetadata** has exclusive time percentage of 2.0%.

TIL: Oh My! Bottlenecks!

- We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

```
// Hardware intrinsics are recognized by name.  
const char* namespaceName = NULL;  
const char* className = NULL;  
const char* methodName = m_interpCeeInfo.getMethodNameFromMetadata((CORINFO_METHOD_HANDLE)methToCall, &className, &namespaceName, NULL);  
if (  
#if defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.X86") == 0 &&  
#elif defined(TARGET_ARM64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.Arm") == 0 &&  
#endif // defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(methodName, "get_IsSupported") == 0  
)  
{  
    GCX_COOP();  
    DoGetIsSupported();  
    didIntrinsic = true;  
}
```

- **TIL:** Obtaining names from metadata is **expensive**.
- `CEEInfo::getMethodNameFromMetadata` has exclusive time percentage of 2.0%.
- **Proposal #1:** Call `getMethodNameFromMetadata` only once for all special cases.

TIL: Oh My! Bottlenecks!

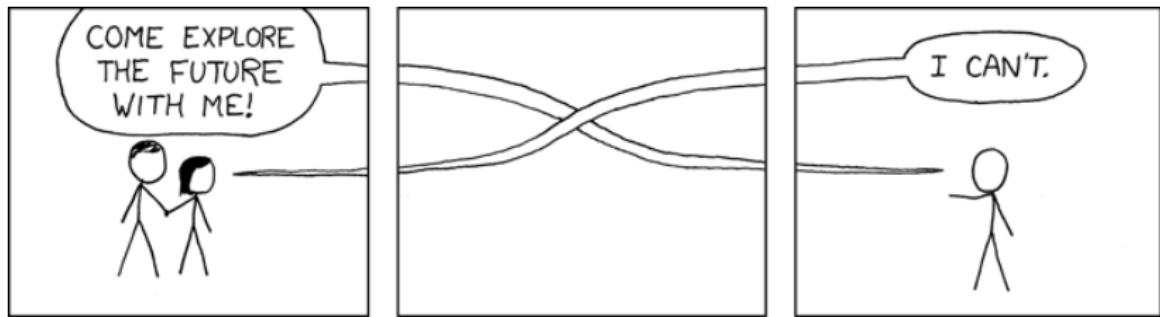
- ▶ We use the names from the metadata to recognize whether a called method is a hardware intrinsic or SIMD method.

```
// Hardware intrinsics are recognized by name.  
const char* namespaceName = NULL;  
const char* className = NULL;  
const char* methodName = m_interpCeeInfo.getMethodNameFromMetadata((CORINFO_METHOD_HANDLE)methToCall, &className, &namespaceName, NULL);  
if (  
#if defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.X86") == 0 &&  
#elif defined(TARGET_ARM64)  
    strcmp(namespaceName, "System.Runtime.Intrinsics.Arm") == 0 &&  
#endif // defined(TARGET_X86) || defined(TARGET_AMD64)  
    strcmp(methodName, "get_IsSupported") == 0  
)  
{  
    GCX_COOP();  
    DoGetIsSupported();  
    didIntrinsic = true;  
}
```

- ▶ **TIL:** Obtaining names from metadata is **expensive**.
- ▶ `CEEInfo::getMethodNameFromMetadata` has exclusive time percentage of 2.0%.
- ▶ **Proposal #1:** Call `getMethodNameFromMetadata` only once for all special cases.
- ▶ **Proposal #2:** Cache the special casing information somewhere (e.g. `CORINFO_INTRINSIC_HWINTRINSIC`).

Where Are We?

- ▶ Why Is the Interpreter So Interesting?
- ▶ ... and Fixes, and Fixes, and Fixes, ...
- ▶ TIL from Beautiful Plots
- ▶ Oh My! Bottlenecks!
- ▶ Where Do We Go from Here? ←



TIL: Where Do We Go From Here?

TIL: Where Do We Go From Here?

- ▶ There are still many **functionality gaps** from the current state of the interpreter to a shippable shape.

TIL: Where Do We Go From Here?

- ▶ There are still many **functionality gaps** from the current state of the interpreter to a shippable shape.
- ▶ **Hardware intrinsics**

TIL: Where Do We Go From Here?

- ▶ There are still many **functionality gaps** from the current state of the interpreter to a shippable shape.
- ▶ **Hardware intrinsics**
- ▶ **ASP.NET:** Enables the performance measurements (throughput and start-up) of a web application as in the Tiered Compilation experiments

TIL: Where Do We Go From Here?

- ▶ Several broad, possibly interesting **performance** questions:

TIL: Where Do We Go From Here?

- ▶ Several broad, possibly interesting **performance** questions:
- ▶ Background tiering-up requests go into a queue, possibly building up a huge backlog with calls to various functions.
Can we turn this into a **priority queue**? What determines the priority of an instance of tiering-up?

TIL: Where Do We Go From Here?

- ▶ Several broad, possibly interesting **performance** questions:
- ▶ Background tiering-up requests go into a queue, possibly building up a huge backlog with calls to various functions.
Can we turn this into a **priority queue**? What determines the priority of an instance of tiering-up?
- ▶ **3-Level Tiered Compilation** (Interpreter-QuickJIT-Tier1)?

TIL: Where Do We Go From Here?

- ▶ Several broad, possibly interesting **performance** questions:
- ▶ Background tiering-up requests go into a queue, possibly building up a huge backlog with calls to various functions.
Can we turn this into a **priority queue**? What determines the priority of an instance of tiering-up?
- ▶ **3-Level Tiered Compilation** (Interpreter-QuickJIT-Tier1)?
- ▶ **Synchronous** tiering-up?

TIL: Where Do We Go From Here?

- ▶ Several broad, possibly interesting **performance** questions:
- ▶ Background tiering-up requests go into a queue, possibly building up a huge backlog with calls to various functions.
Can we turn this into a **priority queue**? What determines the priority of an instance of tiering-up?
- ▶ **3-Level Tiered Compilation** (Interpreter-QuickJIT-Tier1)?
- ▶ **Synchronous** tiering-up?
- ▶ **On Stack Replacement**?

TIL: This Internship Is Great! ❤

TIL: This Internship Is Great! ❤



Questions & Answers

Appendix

- ▶ **Issue #1:** The C++ compiler treats `reinterpret_cast` from a pointer to an integral type as an error.

Appendix

- ▶ **Issue #1:** The C++ compiler treats `reinterpret_cast` from a pointer to an integral type as an error.
- ▶ **Fix:**

```
reinterpret_cast<short>(ArgSlotEndianessFixup(directOffset, sizeof(void*)));
static_cast<short>(reinterpret_cast<intptr_t>(ArgSlotEndianessFixup(directOffset, sizeof(void*))));
```

Appendix

- ▶ **Issue #2:** The C++ compiler treats the conversion from a generic type T to an integral type as an error.

Appendix

- ▶ **Issue #2:** The C++ compiler treats the conversion from a generic type T to an integral type as an error.
- ▶ **Fix:** There are finitely many integral types as CORINFO_TYPE. Hence we can condition on `sizeof(T)` to determine which integral type to which T is equivalent.

```
// Widen narrow types.
int ires;
switch (sizeof(T))
{
    case 1:
        ires = std::is_same<T, INT8>::value ?
            static_cast<int>(reinterpret_cast<INT8*>(res_ptr)[index]) :
            static_cast<int>(reinterpret_cast<UINT8*>(res_ptr)[index]);
        break;
    case 2:
        ires = std::is_same<T, INT16>::value ?
            static_cast<int>(reinterpret_cast<INT16*>(res_ptr)[index]) :
            static_cast<int>(reinterpret_cast<UINT16*>(res_ptr)[index]);
        break;
    case 4:
        ires = std::is_same<T, INT32>::value ?
            static_cast<int>(reinterpret_cast<INT32*>(res_ptr)[index]) :
            static_cast<int>(reinterpret_cast<UINT32*>(res_ptr)[index]);
        break;
    default:
        _ASSERTE_MSG(false, "This should have exhausted all the possible sizes.");
        break;
}
```

Appendix

- ▶ **Issue #3:** The interpreter does not check the optimization tier before it triggers the **asynchronous** tiering-up. The tiering-up might be triggered again before the optimization tier is **actually** promoted.

Appendix

- ▶ **Issue #3:** The interpreter does not check the optimization tier before it triggers the **asynchronous** tiering-up. The tiering-up might be triggered again before the optimization tier is **actually** promoted.
- ▶ **Issue #4:** LockHolder is missing for synchronization.

Appendix

- ▶ **Issue #3:** The interpreter does not check the optimization tier before it triggers the **asynchronous** tiering-up. The tiering-up might be triggered again before the optimization tier is **actually** promoted.
- ▶ **Issue #4:** LockHolder is missing for synchronization.
- ▶ **Fix:**

```
#ifdef FEATURE_TIERED_COMPILATION
    CodeVersionManager::LockHolder _lockHolder;
    NativeCodeVersion activeCodeVersion = md->GetCodeVersionManager()->GetActiveILCodeVersion(md).GetActiveNativeCodeVersion(md);
    ILCodeVersion ilCodeVersion = activeCodeVersion.GetILCodeVersion();
    if (activeCodeVersion.GetOptimizationTier() == NativeCodeVersion::OptimizationTier0 &&
        !ilCodeVersion.HasAnyOptimizedNativeCodeVersion(activeCodeVersion))
    {
        tieredCompilationManager->AsyncPromoteToTier1(activeCodeVersion, &scheduleTieringBackgroundWork);
    }
#else
#error FEATURE_INTERPRETER depends on FEATURE_TIERED_COMPILATION now
#endif
```

Appendix

- ▶ **Issue #5:** “Bad IL Format” for virtual calls.
- ▶ **Fix:** Enable the CORINFO_CALLINFO_CALLVIRT flag in the call info for virtual calls.

```
combine(CORINFO_CALLINFO_SECURITYCHECKS, CORINFO_CALLINFO_LDFTN),  
&callInfo);  
    combine(CORINFO_CALLINFO_CALLVIRT,  
            combine(CORINFO_CALLINFO_SECURITYCHECKS,  
                    CORINFO_CALLINFO_LDFTN)),  
&callInfo);
```

Appendix

- ▶ **Issue #6:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

Appendix

- ▶ **Issue #6:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
        bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call      bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

Appendix

- ▶ **Issue #6:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
        bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call      bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

- ▶ **“Fix” (a.k.a. Workaround):** Force the interpreter to return `false` for `get_IsSupported()` to disable hardware intrinsics.

Appendix

- ▶ **Issue #6:** The IL implementations of `get_IsSupported()` methods for hardware intrinsics are non-terminating recursions and will be handled by jitted assembly code. The interpreter is unaware of this special case.

```
.method public hidebysig specialname static
        bool  get_IsSupported() cil managed
{
    // Code size      6 (0x6)
    .maxstack  8
    IL_0000:  call      bool System.Runtime.Intrinsics.X86.Sse2::get_IsSupported()
    IL_0005:  ret
} // end of method Sse2::get_IsSupported
```

- ▶ **“Fix” (a.k.a. Workaround):** Force the interpreter to return `false` for `get_IsSupported()` to disable hardware intrinsics.
- ▶ Controlled by the COM+ environment variable
`InterpreterHWIntrinsicsIsSupportedFalse`

Appendix

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

Appendix

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

Appendix

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies.

Appendix

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies.
- ▶ **Fix #2:** Implement hardware intrinsics methods in the interpreter.

Appendix

- ▶ **Issue of “Fix”:** Arbitrary combinations of different code generators (JIT, CrossGen, interpreter, etc.) in the same process might cause **inconsistent** values of IsSupported. (@jkotas)

```
readonly static RenderingEngine s_engine = Avx.IsSupported ? new AvxRenderingEngine() : new GenericRenderingEngine();  
  
class GenericRenderingEngine : RenderingEngine  
{  
    virtual void Render()  
    {  
        Debug.Assert(!Avx.IsSupported);  
        // ...  
    }  
  
    // ...  
}
```

- ▶ **Fix #1:** Depend on Fallback to JIT for hardware intrinsics methods and let JIT generate valid method bodies. ↗
- ▶ **Fix #2:** Implement hardware intrinsics methods in the interpreter.